

Policy Evaluation

Objectives

- Understand the distinction between policy evaluation and control
- Explain the setting in which dynamic programming can be applied, as well as its limitations
- Apply iterative policy evaluation to compute value functions

Policy Evaluation vs. Control

- Policy evaluation is the task of determining the value function for a specific policy.
- Control is the task of finding a policy to obtain as much reward as possible: finding a policy which maximizes the value function.

Policy Evaluation vs. Control

- Imagine someone hands you a policy and your job is to determine how good that policy is. Policy evaluation is the task of determining the state value function v_π for a particular policy π .

$$\pi \longrightarrow v_\pi$$

Recall that

$$v_\pi(s) \doteq \mathbb{E}_\pi [G_t | S_t = s]$$

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Policy Evaluation vs. Control

- The Bellman equation reduces the problem of finding v_π to a system of linear equations
- The problem of policy evaluation reduces to solving this system of linear equations.
- The iterative solution methods of dynamic programming are more suitable

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_\pi(s')]$$

Recall that

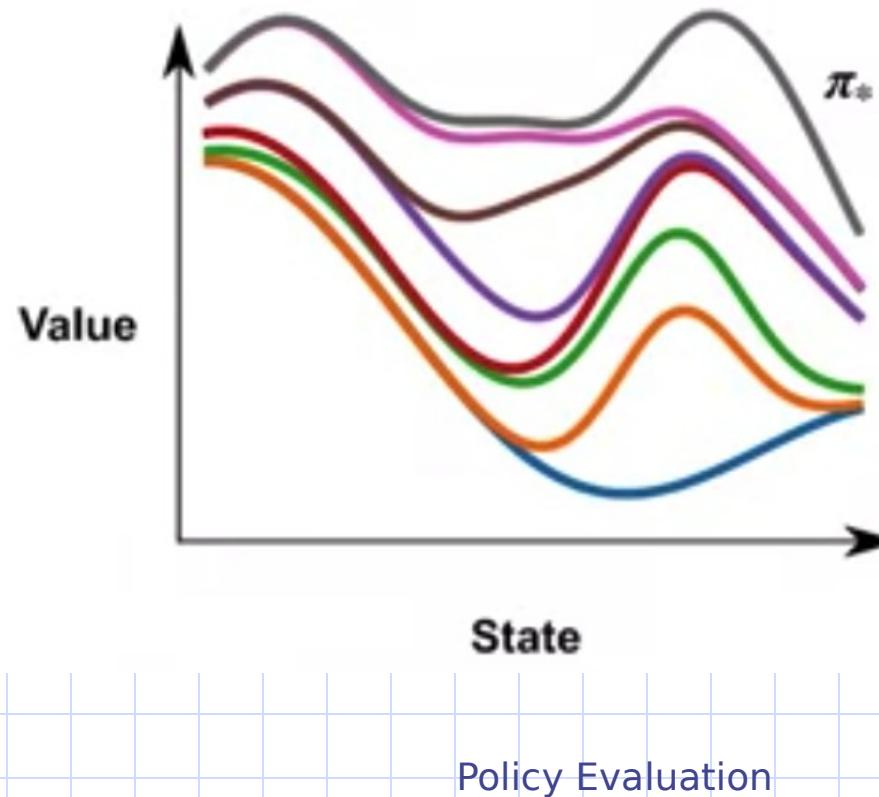


In practice



Policy Evaluation vs. Control

- Control is the task of improving a policy..



Policy Evaluation vs. Control

- The dynamics of the environment: p
- We use dynamic programming methods to compute value functions and optimal policies given a model of the MDP..



Iterative Policy Evaluation

- The Bellman equation gives us a recursive expression for $V \pi$.
- Iteratively refine our estimate of the value function.

$$v_{\pi}(s) = \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_{\pi}(s')]$$

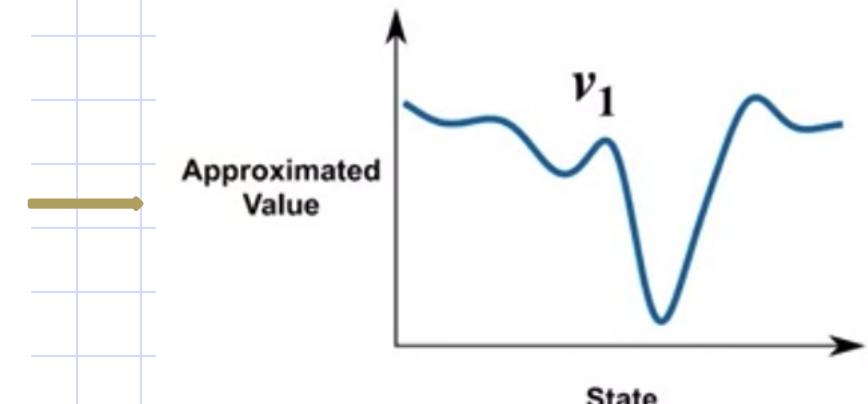
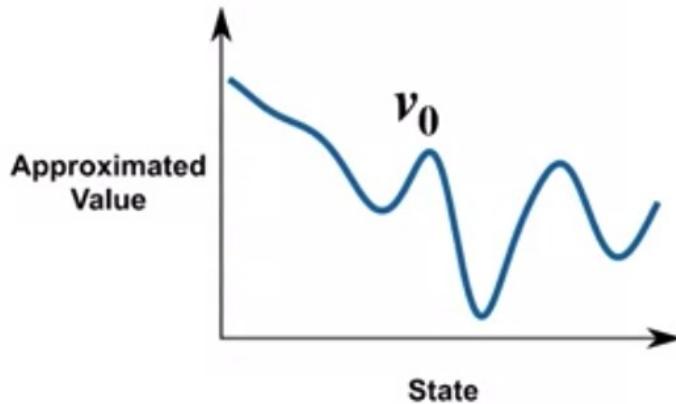


$$v_{k+1}(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_k(s')]$$

Iterative Policy Evaluation

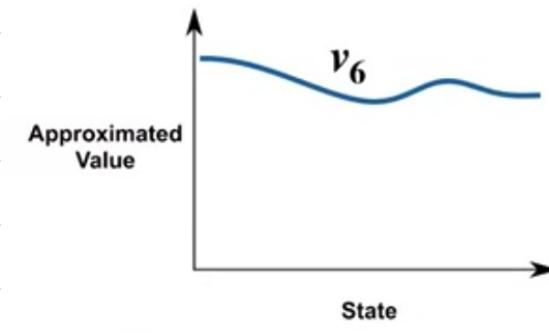
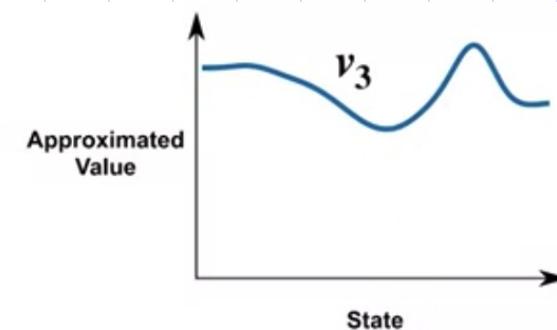
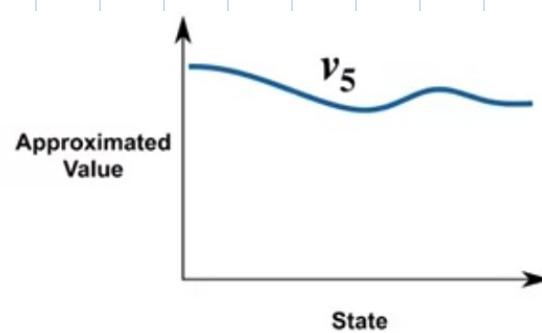
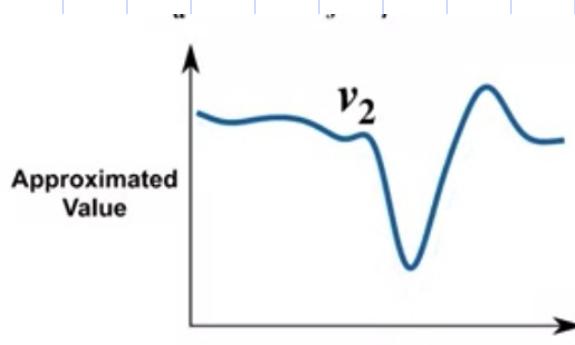
- An arbitrary initialization for our approximate value function- v_0
- Using update rule:

$$v_{k+1}(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_k(s')]$$



Iterative Policy Evaluation

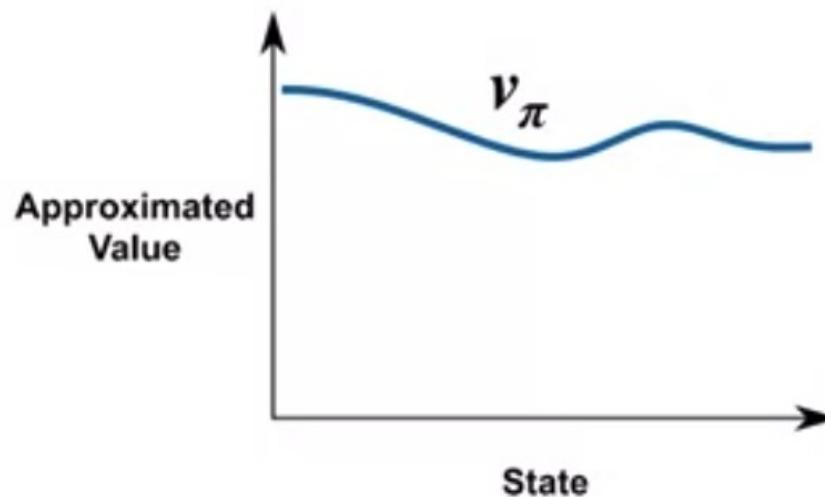
Update V using interactive



Iterative Policy Evaluation

- For any choice of v_0 , v_k will converge to v_π in the limit as k approaches infinity.

$$v_k(s) = \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_k(s')]$$



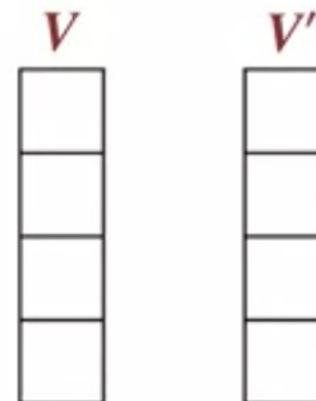
For any v_0

$$\lim_{k \rightarrow \infty} v_k = v_\pi$$

Iterative Policy Evaluation

- To implement iterative policy evaluation, we store two arrays.
 - One array- V stores the current approximate value function.
 - Another array, V' , stores the updated values.

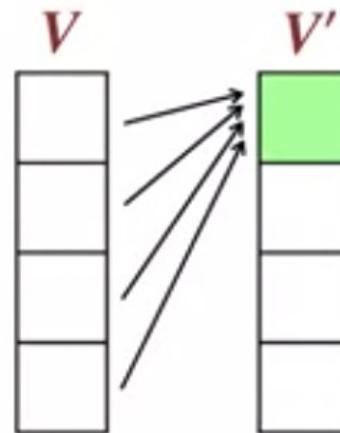
$$V'(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$



Iterative Policy Evaluation

- By using two arrays, we can compute the new values from the old one state at a time without the old values being changed in the process.

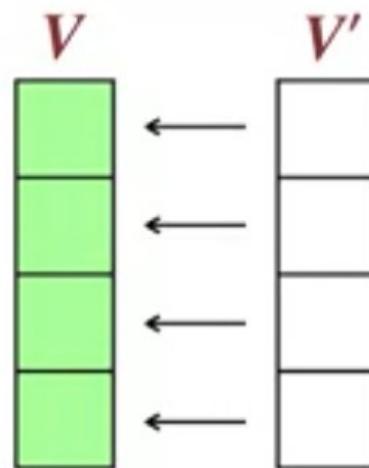
$$\mathbf{V}'(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma \mathbf{V}(s')]$$



Iterative Policy Evaluation

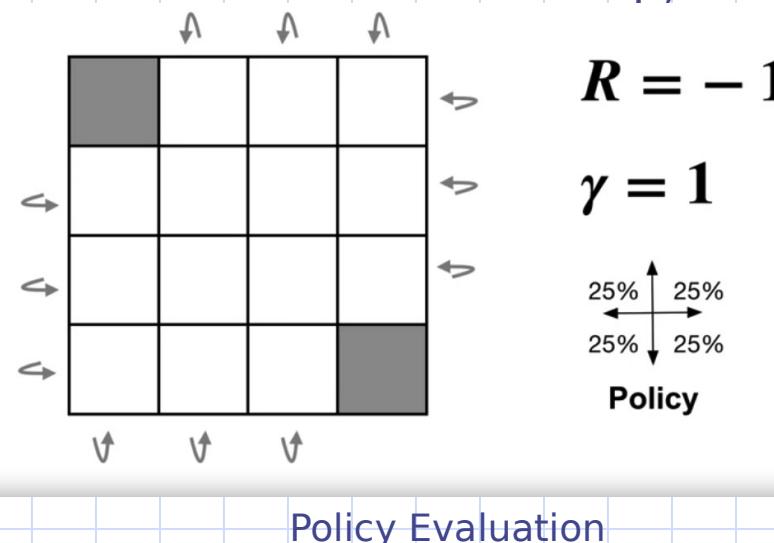
- At the end of a full sweep, we can write all the new values into V ; then we do the next iteration..

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$



Iterative Policy Evaluation

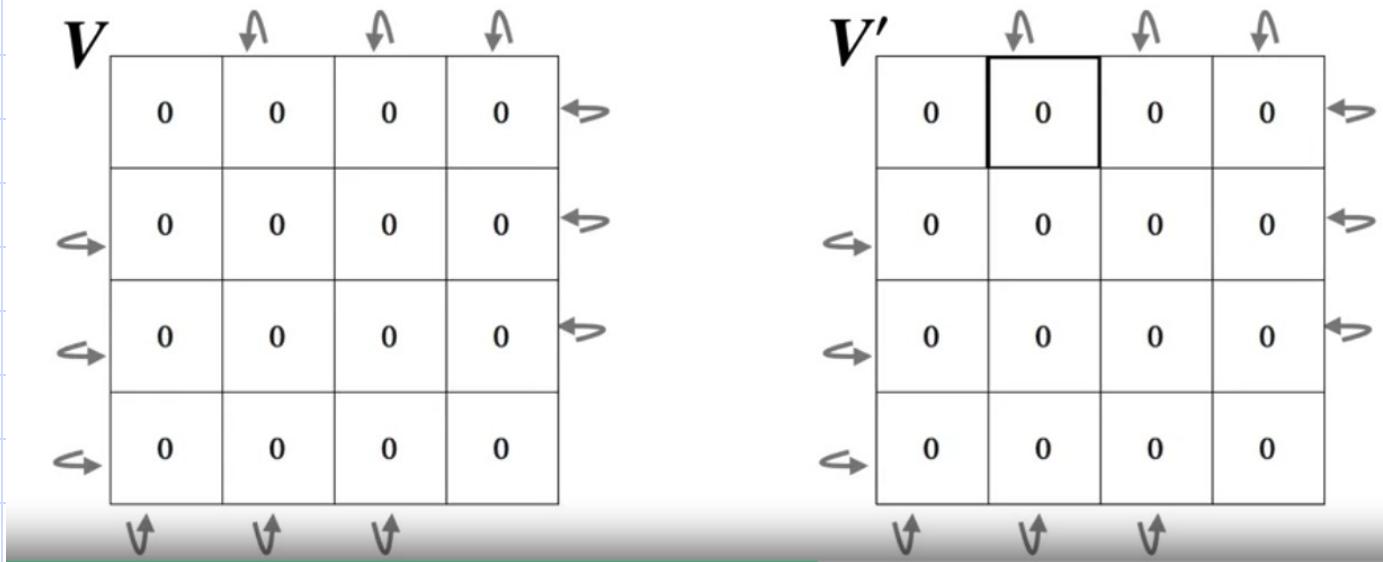
- The four-by-four grid world
- The terminal state located in the top left and bottom right corners.
- The reward will be minus one for every transition.
- gamma equals 1
- Four possible actions in each state up, down, left, and right.



Iterative Policy Evaluation

□ Calculate:

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$



Iterative Policy Evaluation

□ Calculate:

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

$$0.25 * (-1 + 0) + 0.25 * (-1 + 0) + 0.25 * (-1 + 0) + 0.25 * (-1 + 0) = -1$$

V	↓	↓	↓	↔
←	0	0	0	↔
←	0	0	0	↔
←	0	0	0	↔
←	0	0	0	↔
↑	↑	↑	↑	

V'	↓	↓	↓	↔
←	0	0	0	↔
←	0	0	0	↔
←	0	0	0	↔
←	0	0	0	↔
↑	↑	↑	↑	

Iterative Policy Evaluation

□ Calculate:

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

$$0.25 * (-1 + 0) + 0.25 * (-1 + 0) + 0.25 * (-1 + 0) + 0.25 * (-1 + 0) = -1$$

V				
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

V'				
0	-1	-1	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Iterative Policy Evaluation

□ Calculate:

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

V

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

↓ ↓ ↓

↔ ↔ ↔

V'

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

↓ ↓ ↓

↔ ↔ ↔

Iterative Policy Evaluation

- After completing this full sweep, we copy the updated values from V' prime to V :

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

V				
0	-1	-1	-1	↔
-1	-1	-1	-1	↔
-1	-1	-1	-1	↔
-1	-1	-1	0	↑

V'				
0	-1	-1	-1	↔
-1	-1	-1	-1	↔
-1	-1	-1	-1	↔
-1	-1	-1	0	↑

Iterative Policy Evaluation

- Iterative Policy Evaluation, for estimating $V \sim V_{\pi}$:

Input π , the policy to be evaluated

$$V \leftarrow \vec{0}, V' \leftarrow \vec{0}$$

Loop:

$$\Delta \leftarrow 0$$

Loop for each $s \in \mathcal{S}$:

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s',r} p(s',r | s,a) [r + \gamma V(s')]$$

$$\Delta \leftarrow \max(\Delta, |V'(s) - V(s)|)$$

$$V \leftarrow V'$$

until $\Delta < \theta$ (a small positive number)

Output $V \approx v_{\pi}$

Iterative Policy Evaluation

- Our value of 0.001 for the stopping parameter theta

- $V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$

$$\theta = 0.001 \Delta = 1.0$$

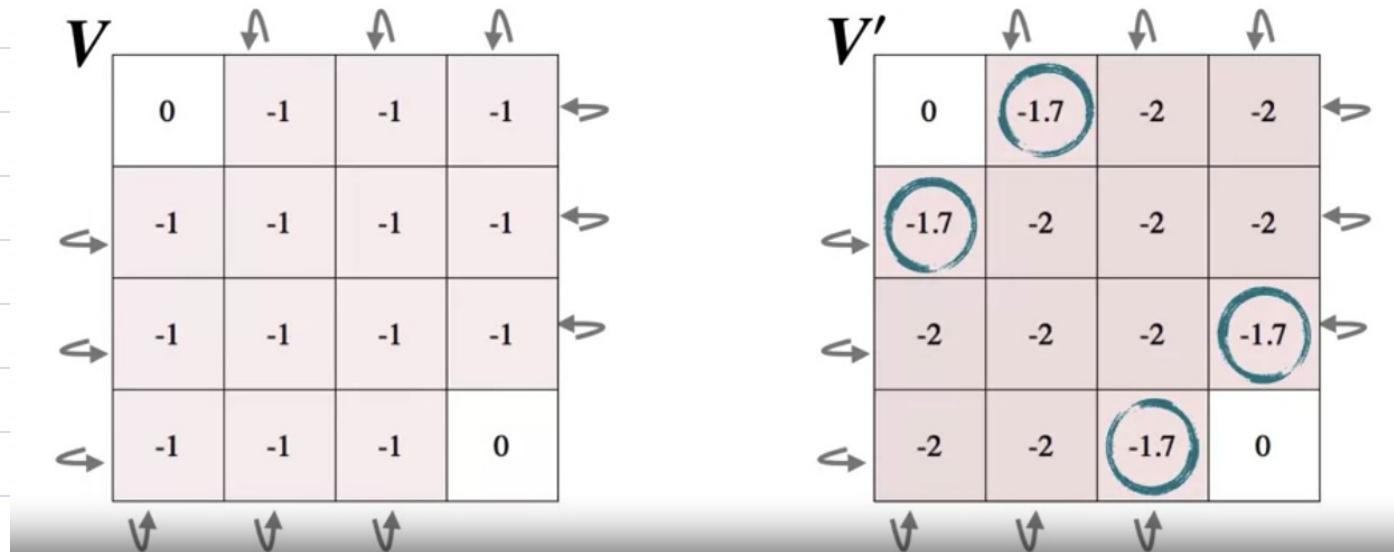
V	↓	↓	↓	↔
↔	0	-1	-1	-1
↔	-1	-1	-1	-1
↔	-1	-1	-1	-1
↔	-1	-1	-1	0
↑	↑	↑	↑	↔

V'	↓	↓	↓	↔
↔	0	-1	-1	-1
↔	-1	-1	-1	-1
↔	-1	-1	-1	-1
↔	-1	-1	-1	0
↑	↑	↑	↑	↔

Iterative Policy Evaluation

- The second sweep

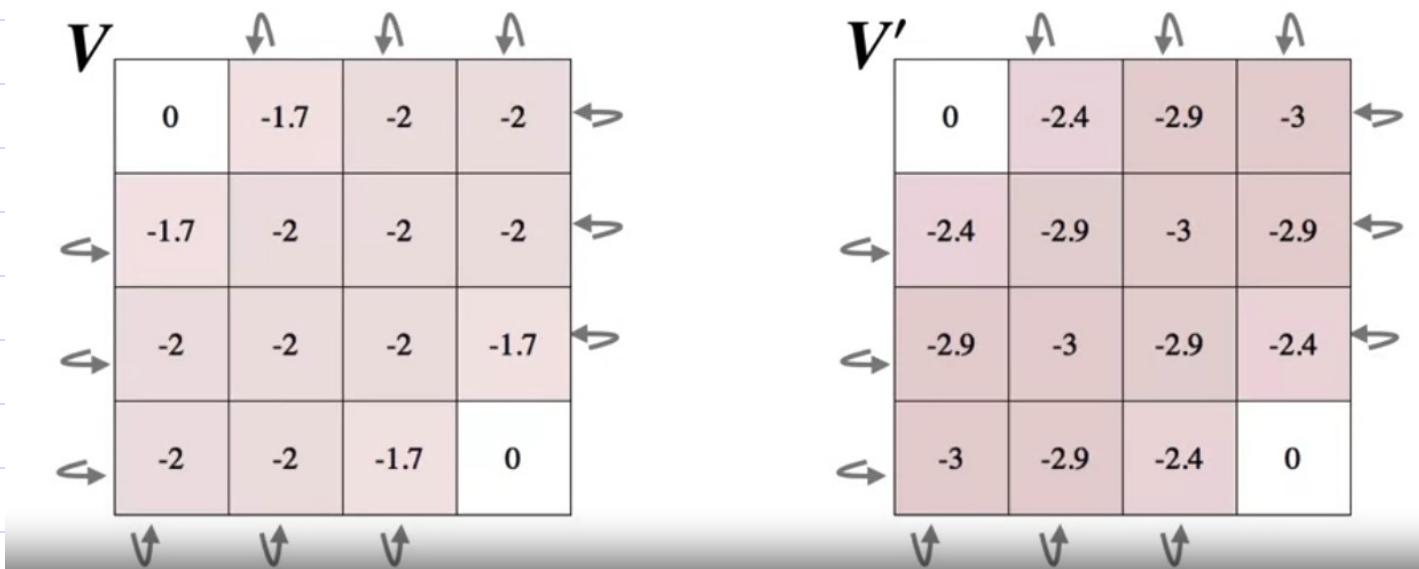
$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$



Iterative Policy Evaluation

- One more sweep

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$



Iterative Policy Evaluation

- More sweep

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

V	↓	↓	↓	↔
0	-2.4	-2.9	-3	↔
↔	-2.4	-2.9	-3	↔
↔	-2.9	-3	-2.9	↔
↔	-3	-2.9	-2.4	↔
↑	↑	↑	↑	

V'	↓	↓	↓	↔
0	-3.1	-3.8	-4	↔
↔	-3.1	-3.7	-3.9	↔
↔	-3.8	-3.9	-3.7	↔
↔	-4	-3.8	-3.1	↔
↑	↑	↑	↑	

Iterative Policy Evaluation

- More sweep

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

V

0	-3.1	-3.8	-4
-3.1	-3.7	-3.9	-3.8
-3.8	-3.9	-3.7	-3.1
-4	-3.8	-3.1	0

↓ ↓ ↓ ↓ ↑ ↑ ↑ ↑

V'

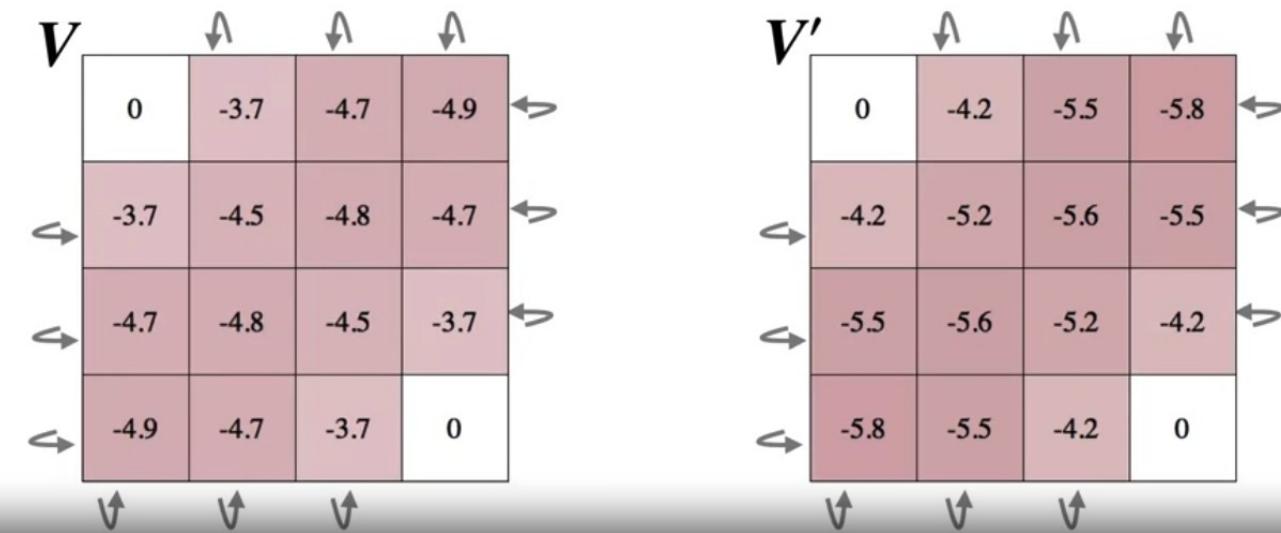
0	-3.7	-4.7	-4.9
-3.7	-4.5	-4.8	-4.7
-4.7	-4.8	-4.5	-3.7
-4.9	-4.7	-3.7	0

↓ ↓ ↓ ↓ ↑ ↑ ↑ ↑

Iterative Policy Evaluation

- More sweep

$$V'(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$



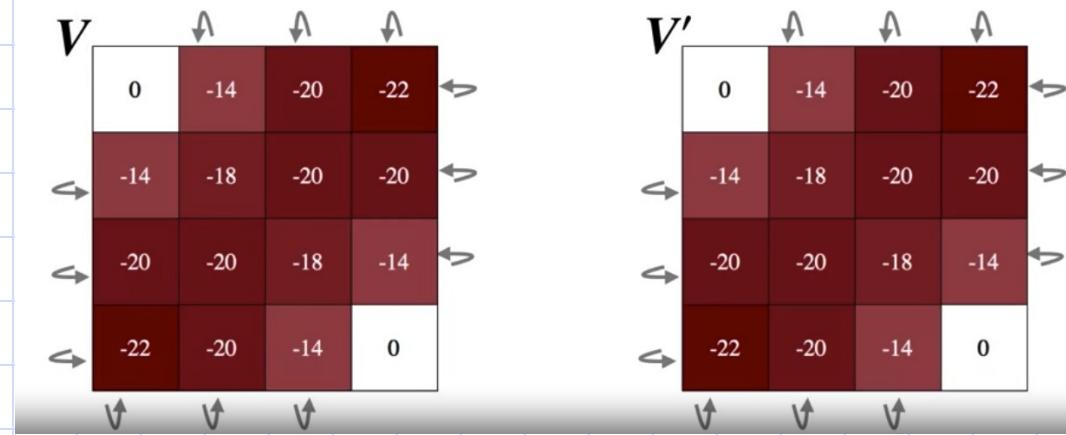
Iterative Policy Evaluation

- Keep running until our maximum delta is less than

theta.

$$V'(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma V(s')]$$

$$\Delta < 0.001$$



- Our approximate value function has converged to the value function for the random policy

Summary

- Understand the distinction between policy evaluation and control
- Explain the setting in which dynamic programming can be applied, as well as its limitations
- Apply iterative policy evaluation to compute value functions

Q & A