

Actor-Critic for Continuing Tasks

Objectives

- Derive a sample-based estimate for the gradient
- Understand the Actor-Critic algorithm

Estimating the Policy Gradient

Definition:

- Estimating the policy gradient is a crucial step in policy gradient algorithms
- It allows us to update the policy parameters in the direction that maximizes the expected return.

Estimating the Policy Gradient

- The gradient of the average reward:

$$\nabla r(\pi) = \sum_s \mu(s) \sum_a \nabla \pi(a | s, \theta) q_\pi(s, a)$$

- We simply make updates from states we observe while following policy π .

$S_0, A_0, R_1, S_1, A_1, \dots, S_t, A_t, R_{t+1}, \dots$

Estimating the Policy Gradient

- This gradient from state S_t provides an approximation to the gradient of the average reward.
- The stochastic gradient descent update looks like

$$\text{for the } \nabla r(\pi) = \sum_s \mu(s) \sum_a \nabla \pi(a | s, \theta) q_\pi(s, a)$$

$$\theta_{t+1} \doteq \theta_t + \alpha \sum_a \nabla \pi(a | S_t, \theta_t) q_\pi(S_t, a)$$

$S_0, A_0, R_1, S_1, A_1, \dots, S_t, A_t, R_{t+1}, \dots$

Actor-Critic for Continuing Tasks

Estimating the Policy Gradient

- Estimating the policy gradient methods
 - Finite Difference Methods: These methods approximate the gradient by perturbing the policy parameters and observing the change in the expected return.
 - Score Function Methods: Score function methods, directly estimate the gradient of the expected return using samples obtained from interactions with the environment.

Estimating the Policy Gradient

- Estimating the policy gradient methods
 - Likelihood Ratio Methods: Likelihood ratio methods, further improve the efficiency of gradient estimation by subtracting a baseline function from the returns before computing the gradient.
 - Actor-Critic Methods: Actor-critic methods combine policy gradient estimation with value function estimation.

Estimating the Policy Gradient

- Estimating the policy gradient methods
 - Natural Policy Gradient: Natural policy gradient methods incorporate the geometry of the parameter space into gradient estimation to improve convergence properties and stability.

Actor-Critic Algorithm

- The Actor-Critic algorithm combines elements of both value-based and policy-based methods.
- It consists of two main components: the actor and the critic:
 - **Actor:** The actor is responsible for learning the policy
 - **Critic:** The critic is responsible for evaluating the policy.
- The Actor-Critic algorithm operates by iteratively updating the actor and critic networks based on observed experiences in the environment.

Actor-Critic Algorithm

Actor- Critic Algorithm:

- Step 1- Initialization: Initialize the actor and critic networks with random parameters.
- Step 2- Interact with Environment: Sample trajectories by following the current policy in the environment.
- Step 3- Compute Returns: Compute the returns for each time step in the trajectory.
- Step 4- Update Critic: Use the returns to update the parameters of the critic network

Actor-Critic Algorithm

Actor- Critic Algorithm:

- Step 5- Compute Advantages: Compute advantages for each time step by subtracting the estimated value from the observed return.
- Step 6- Update Actor: Update the parameters of the actor network using policy gradients
- Step 7- Repeat: Repeat steps 2-6 for multiple episodes or until convergence criteria are met.

Actor-Critic Algorithm

- The Actor-Critic algorithm in a simple grid world environment

```
1 # 3.11 Actor-Critic Algorithm- HoaDNT@fe.edu.vn
2 import numpy as np
3
4 class GridWorld:
5     def __init__(self):
6         self.grid_size = (3, 3)
7         self.num_actions = 4 # Up, Down, Left, Right
8         self.start_state = (0, 0)
9         self.goal_state = (2, 2)
10
11    def step(self, state, action):
12        # Define the dynamics of the environment
13        row, col = state
14        if action == 0: # Up
15            row = max(0, row - 1)
16        elif action == 1: # Down
17            row = min(self.grid_size[0] - 1, row + 1)
18        elif action == 2: # Left
19            col = max(0, col - 1)
20        elif action == 3: # Right
21            col = min(self.grid_size[1] - 1, col + 1)
22        next_state = (row, col)
23        reward = 0
24        if next_state == self.goal_state:
25            reward = 1 # Reward of +1 upon reaching the goal state
26        return next_state, reward
27
```

Actor-Critic Algorithm

```
28 class ActorCritic:
29     def __init__(self, num_actions, alpha_actor, alpha_critic, gamma):
30         self.num_actions = num_actions
31         self.alpha_actor = alpha_actor
32         self.alpha_critic = alpha_critic
33         self.gamma = gamma
34         self.actor_params = np.zeros((3, 3, num_actions)) # Tabular actor parameters
35         self.critic_values = np.zeros((3, 3)) # Tabular critic values
36
37     def select_action(self, state):
38         # Select action probabilistically based on actor parameters
39         action_probs = self.softmax(self.actor_params[state])
40         action = np.random.choice(self.num_actions, p=action_probs)
41         return action
42
43     def update(self, state, action, reward, next_state):
44         # Compute TD error (advantage)
45         td_error = reward + self.gamma * self.critic_values[next_state] - self.critic_values[state]
46
47         # Update critic values
48         self.critic_values[state] += self.alpha_critic * td_error
49
50         # Update actor parameters
51         self.actor_params[state][action] += self.alpha_actor * td_error
52
53     def softmax(self, x):
54         e_x = np.exp(x - np.max(x))
55         return e_x / e_x.sum(axis=0)
```

Actor-Critic Algorithm

```
57 # Create a grid world environment
58 grid_world = GridWorld()
59
60 # Create an Actor-Critic agent
61 num_actions = 4 # Up, Down, Left, Right
62 alpha_actor = 0.1
63 alpha_critic = 0.1
64 gamma = 0.9
65 actor_critic_agent = ActorCritic(num_actions, alpha_actor, alpha_critic, gamma)
66
67 # Train the Actor-Critic agent
68 num_episodes = 1000
69 for _ in range(num_episodes):
70     state = grid_world.start_state
71     while state != grid_world.goal_state:
72         action = actor_critic_agent.select_action(state)
73         next_state, reward = grid_world.step(state, action)
74         actor_critic_agent.update(state, action, reward, next_state)
75         state = next_state
76
77 # Evaluate the learned policy
78 total_reward = 0
79 state = grid_world.start_state
80 while state != grid_world.goal_state:
81     action = actor_critic_agent.select_action(state)
82     next_state, reward = grid_world.step(state, action)
83     total_reward += reward
84     state = next_state
85
86 print("Total reward obtained by learned policy:", total_reward)
```

Actor-Critic Algorithm

- Run that code and show a total reward obtained by learned policy

Summary

- Derive a sample-based estimate for the gradient
- Understand the Actor-Critic algorithm

Q & A